

Optimizing Loan Limit Increases: A Data-Driven Approach

Notebook Summary (Introduction Section):

This notebook explores an optimal strategy for loan limit increases using machine learning, operations research, and stochastic modeling. The dataset consists of 30,000 customer loan records, including repayment history, eligibility criteria, and default risk.

Objectives:

1. Perform Exploratory Data Analysis (EDA) to understand borrower behavior.
2. Develop a forecasting model for customer uptake of loan increases.
3. Implement an optimization model to maximize profitability while controlling risk.
4. Use probabilistic techniques (Markov chains, Monte Carlo simulations) to model borrower transitions
5. Simulate loan lifecycle scenarios and evaluate strategies under different economic conditions.
6. Provide recommendations for real-world implementation.

Step 1: Understanding the Dataset

Before we jump into modeling, we need to understand the structure of your dataset. Upload the dataset, and we'll start with Exploratory Data Analysis (EDA) to examine the following:

1. Column names and data types
2. Missing values and outliers
3. Summary statistics
4. Key trends in loan amounts, repayment behavior, and risk categories

Read the Xl

```
!pip install pandas openpyxl
```

```
Requirement already satisfied: pandas in c:\users\admin\anaconda3\lib\site-packages (2.2.2)
```

```
Requirement already satisfied: openpyxl in c:\users\admin\anaconda3\lib\site-packages (3.1.5)
```

```
Requirement already satisfied: numpy>=1.26.0 in c:\users\admin\anaconda3\lib\site-packages (from pandas) (1.26.4)
```

```
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\admin\anaconda3\lib\site-packages (from pandas) (2.9.0.post0)
```

```
Requirement already satisfied: pytz>=2020.1 in c:\users\admin\anaconda3\lib\site-packages (from pandas) (2024.1)
```

```
Requirement already satisfied: tzdata>=2022.7 in c:\users\admin\
```

```
anaconda3\lib\site-packages (from pandas) (2023.3)
Requirement already satisfied: et-xmlfile in c:\users\admin\anaconda3\
lib\site-packages (from openpyxl) (1.1.0)
Requirement already satisfied: six>=1.5 in c:\users\admin\anaconda3\
lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)
```

Import main libraries:

```
import pandas as pd

# Load the DataSet
df = pd.read_excel('loan_limit_increases.xlsx', header=1) # Set
header to the second row (index 1)

# Display the first few rows
df.head()
```

	Customer ID	Initial Loan (\$)	Days Since Last Loan	On-time Payments (%)
0	1001	1360	72	90.41
1	1002	4272	54	90.32
2	1003	3592	242	85.56
3	1004	966	301	95.86
4	1005	4926	352	94.80

	No. of Increases in 2023	Total Profit Contribution (\$)
0	5	120
1	0	0
2	4	80
3	0	0
4	0	0

Explore the dataset:

1. Check for missing values, data types, and basic statistics.
2. Identify key features (e.g., loan amount, repayment status, risk category).

Define problem scope:

3. Understand how loan limits are currently assigned.
4. Identify potential risk factors affecting limit increases.

```
# Basic info
```

```
df.info()
```

```
# Check missing values
```

```
df.isnull().sum()
```

```
# Summary statistics
```

```
df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 30000 entries, 0 to 29999
```

```
Data columns (total 6 columns):
```

#	Column	Non-Null Count	Dtype
0	Customer ID	30000 non-null	int64
1	Initial Loan (\$)	30000 non-null	int64
2	Days Since Last Loan	30000 non-null	int64
3	On-time Payments (%)	30000 non-null	float64
4	No. of Increases in 2023	30000 non-null	int64
5	Total Profit Contribution (\$)	30000 non-null	int64

```
dtypes: float64(1), int64(5)
```

```
memory usage: 1.4 MB
```

	Customer ID	Initial Loan (\$)	Days Since Last Loan	\
count	30000.000000	30000.000000	30000.000000	
mean	16000.500000	2752.532600	181.475900	
std	8660.398374	1293.027316	105.066112	
min	1001.000000	500.000000	0.000000	
25%	8500.750000	1630.000000	90.000000	
50%	16000.500000	2761.000000	182.000000	
75%	23500.250000	3866.000000	272.000000	
max	31000.000000	4999.000000	364.000000	

	On-time Payments (%)	No. of Increases in 2023	\
count	30000.000000	30000.000000	
mean	90.024261	2.237500	
std	5.745515	2.076914	
min	80.000000	0.000000	
25%	85.080000	0.000000	
50%	90.050000	3.000000	
75%	95.010000	4.000000	
max	100.000000	5.000000	

	Total Profit Contribution (\$)
count	30000.000000
mean	44.718667
std	46.632106
min	0.000000
25%	0.000000
50%	40.000000

75%	80.000000
max	120.000000

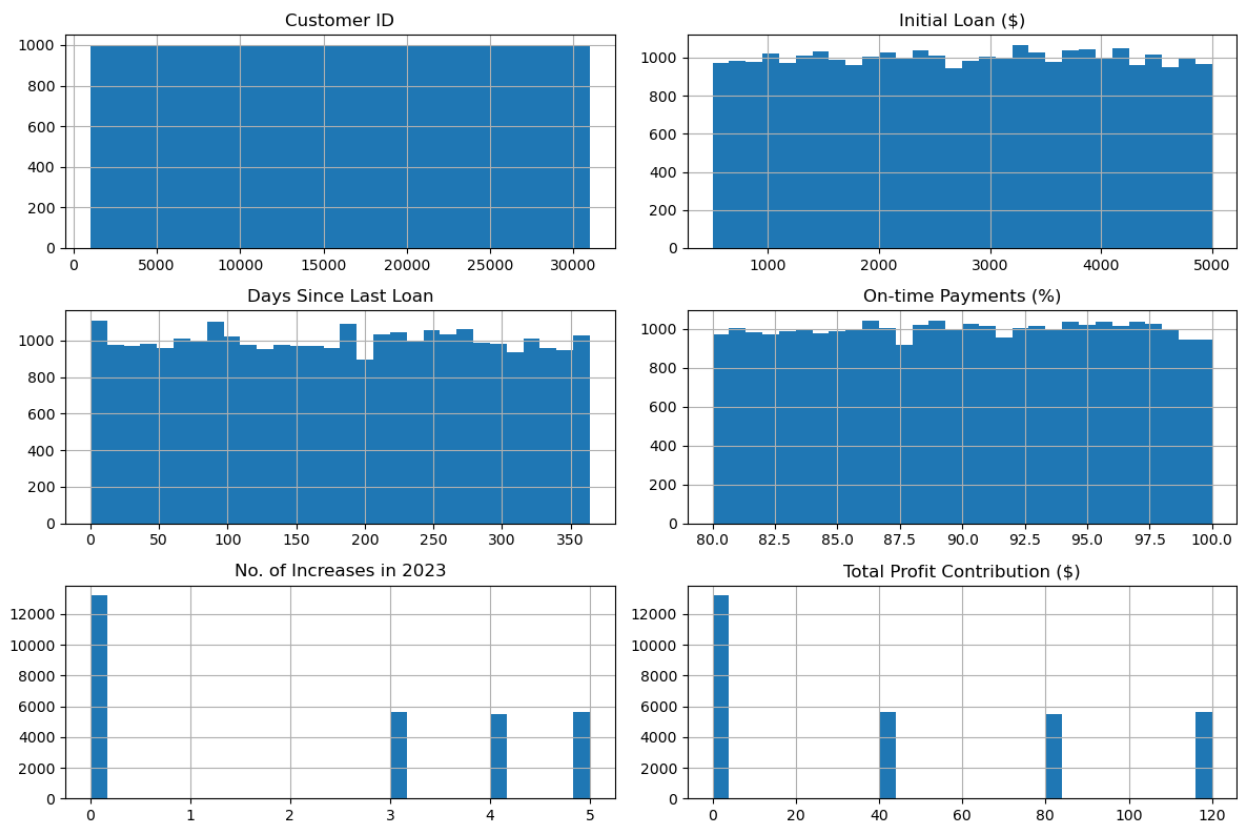
Since there are no missing values, we can now:

1 Understand Data Distribution:

```
import matplotlib.pyplot as plt
import seaborn as sns

# Select numeric columns
numeric_cols = df.select_dtypes(include=['number']).columns

# Plot histograms
df[numeric_cols].hist(figsize = (12,8) , bins = 30)
plt.tight_layout()
plt.show()
```



The fact that we see six different patterns suggests distinct customer behaviors or loan characteristics.

2 Check for Customer Segments:

```
# Check unique risk categories. Let's define risk categories using  
"On-time Payments (%)" as a measure of creditworthiness.
```

```
df.columns
```

```
Index(['Customer ID', 'Initial Loan ($)', 'Days Since Last Loan',  
      'On-time Payments (%)', 'No. of Increases in 2023',  
      'Total Profit Contribution ($)'],  
      dtype='object')
```

```
def assign_risk_category(on_time_payments):  
    if on_time_payments >= 95:  
        return "Prime"  
    elif 80 <= on_time_payments < 95:  
        return "Near-Prime"  
    else:  
        return "Subprime"
```

```
df["risk_category"] = df["On-time Payments  
(%)"].apply(assign_risk_category)
```

```
df.head()
```

	Customer ID	Initial Loan (\$)	Days Since Last Loan	On-time Payments (%) \
0	1001	1360	72	90.41
1	1002	4272	54	90.32
2	1003	3592	242	85.56
3	1004	966	301	95.86
4	1005	4926	352	94.80

	No. of Increases in 2023	Total Profit Contribution (\$)	
0	5	120	Near-Prime
1	0	0	Near-Prime
2	4	80	Near-Prime
3	0	0	Prime

4	0	0	Near-
Prime			

Explore the risk distribution:

1. How many customers fall into each category?
2. Are most customers "Prime" or "Subprime"?

```
df["risk_category"].value_counts()
```

```
risk_category
Near-Prime    22480
Prime         7520
Name: count, dtype: int64
```

Visualize the risk distribution

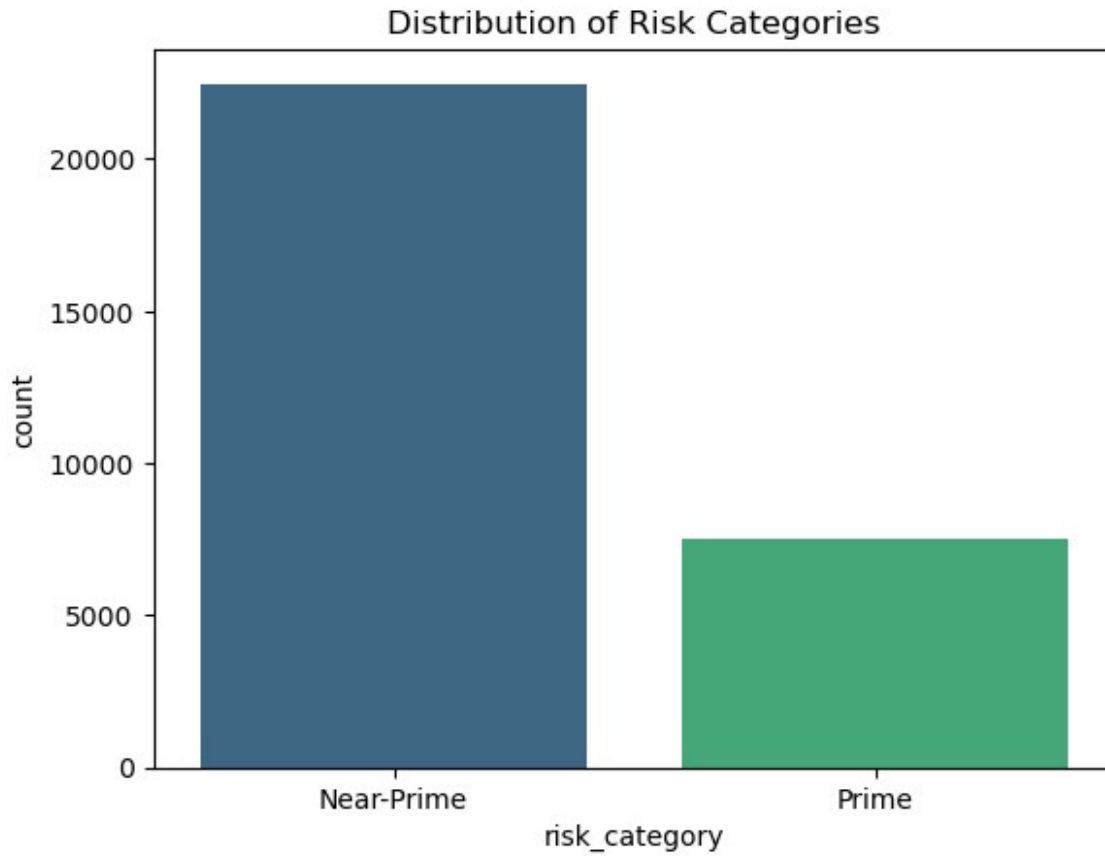
```
import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x = 'risk_category' , data = df , palette = 'viridis')
plt.title('Distribution of Risk Categories')
plt.show()
```

C:\Users\Admin\AppData\Local\Temp\ipykernel_44140\168789355.py:4:
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x = 'risk_category' , data = df , palette = 'viridis')
```



```
## Compare repayment behavior across risk categories
df.groupby("risk_category")["On-time Payments (%)"].mean()

risk_category
Near-Prime    87.540298
Prime         97.449725
Name: On-time Payments (%), dtype: float64
```

Insights:

1. Near-Prime customers are riskier than Prime but still have a relatively high on-time payment rate (~87.5%).
2. The gap (97.45% vs. 87.54%) is significant—Near-Prime customers might require more careful loan limit increases.
3. We need to check if this risk is justified by profit contribution. If Near-Prime customers generate high profits, it might be worth increasing their loan limits despite the slightly higher risk.

```
df.groupby("risk_category")["Total Profit Contribution ($)"].mean()

risk_category
Near-Prime    44.631673
```

```
Prime          44.978723
Name: Total Profit Contribution ($), dtype: float64
```

Key Takeaways:

1. Profit difference is very small (~\$0.35) → Near-Prime customers generate almost the same profit as Prime customers.
2. Near-Prime customers are riskier (lower on-time payments) but profitability is nearly identical to Prime customers.
3. Decision Dilemma: Should we be more lenient with Near-Prime customers if their profit is almost the same as Prime customers?

Note:

If Near-Prime customers default significantly more, we may need stricter loan policies for them. If not, we might be able to expand their loan limits more aggressively

```
# default rates for each category to see if Near-Prime customers
# default more often:

df["Estimated Default Risk (%)"] = 100 - df["On-time Payments (%)"]
df.groupby("risk_category")["Estimated Default Risk (%)"].mean()

risk_category
Near-Prime    12.459702
Prime         2.550275
Name: Estimated Default Risk (%), dtype: float64
```

Key Insights

1. Near-Prime customers are nearly 5x riskier than Prime customers.
2. We must carefully manage loan limit increases for Near-Prime customers.
3. A risk-based loan policy is essential to balance profitability vs. default risk.

```
# Profitability vs. Default Risk

df.groupby("risk_category")["Total Profit Contribution ($)"].mean()

risk_category
Near-Prime    44.631673
Prime         44.978723
Name: Total Profit Contribution ($), dtype: float64
```

Profitability vs. Default Risk Analysis

Key Insights from the Results

1. Near-Prime customers: Average profit contribution \$44.63, with 12.46% default risk.

2. Prime customers: Average profit contribution \$44.98, with 2.55% default risk.

Profit-to-Risk Ratio:

```
df.groupby("risk_category").apply(lambda x: x["Total Profit  
Contribution ($)"].mean() / x['Estimated Default Risk (%)'].mean())
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_44140\3061705377.py:1:  
DeprecationWarning: DataFrameGroupBy.apply operated on the grouping  
columns. This behavior is deprecated, and in a future version of  
pandas the grouping columns will be excluded from the operation.  
Either pass `include_groups=False` to exclude the groupings or  
explicitly select the grouping columns after groupby to silence this  
warning.
```

```
df.groupby("risk_category").apply(lambda x: x["Total Profit  
Contribution ($)"].mean() / x['Estimated Default Risk (%)'].mean())
```

```
risk_category  
Near-Prime    3.582082  
Prime         17.636811  
dtype: float64
```

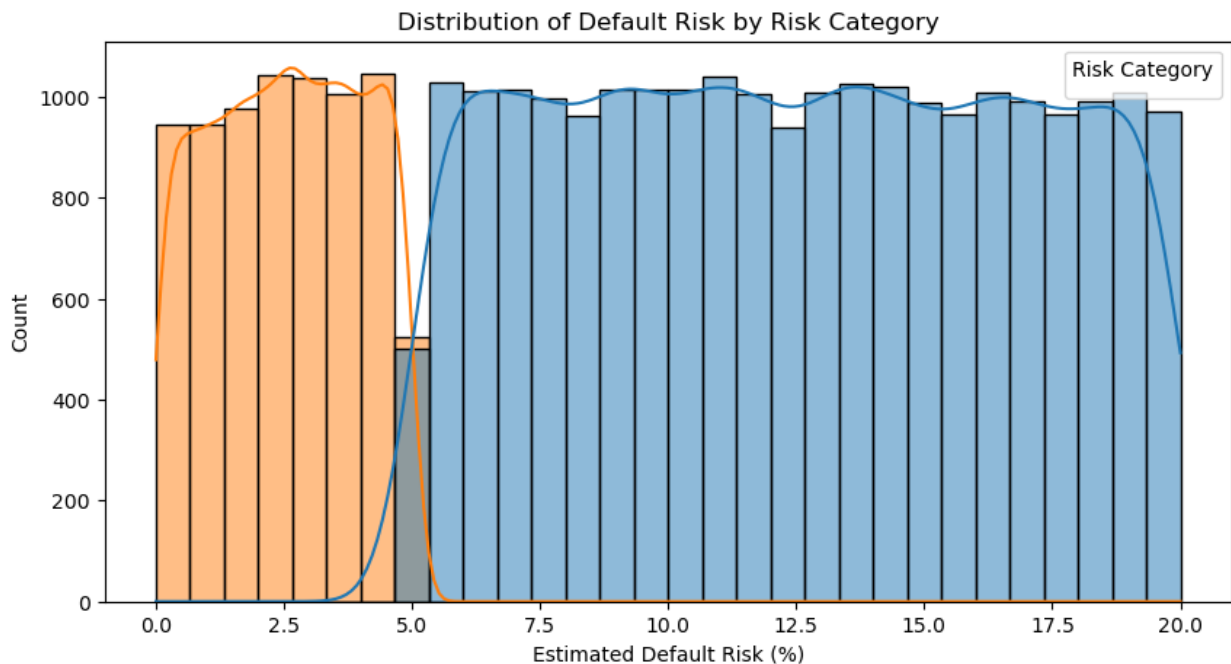
This unusual trend, let's investigate further before making any conclusions

1. Check Distribution of Default Risk in Each Category We'll visualize and compare how default risk is spread across Prime vs. Near-Prime customers.
2. Analyze Loan & Repayment Behavior We'll check if Prime customers take higher loans, have longer repayment periods, or more increases—this could explain the trend.
3. Look for Outliers or Data Issues We'll see if extreme values are skewing our results.

```
import seaborn as sns  
import matplotlib.pyplot as plt  
  
plt.figure(figsize=(10,5))  
sns.histplot(data=df, x='Estimated Default Risk (%)',  
hue='risk_category', kde=True, bins=30)  
plt.title('Distribution of Default Risk by Risk Category')  
plt.xlabel('Estimated Default Risk (%)')  
plt.ylabel('Count')  
plt.legend(title='Risk Category')  
plt.show()
```

```
C:\Users\Admin\AppData\Local\Temp\ipykernel_44140\338225567.py:9:  
UserWarning: No artists with labels found to put in legend. Note that  
artists whose label start with an underscore are ignored when legend()
```

```
is called with no argument.  
plt.legend(title='Risk Category')
```



Insights:

Why do Near-Prime customers have lower estimated risk?

1. There could be a labeling issue in how risk categories were assigned.

1. The model used to estimate default risk might be biased.
2. Prime customers shouldn't have such high default.

2. Their loan amounts & repayment history.

External factors affecting their performance (e.g., economic shifts).

Check Average Loan Amount per Risk Category

Do Prime customers take significantly larger loans, leading to higher estimated risk?

```
df.groupby('risk_category')['Initial Loan ($)'].mean()
```

```
risk_category  
Near-Prime    2746.729537  
Prime         2769.880053  
Name: Initial Loan ($), dtype: float64
```

□ Not a huge difference—both categories borrow almost the same amount.

```
Cell In[127], line 1
    □ Not a huge difference—both categories borrow almost the same
amount.
    ^
SyntaxError: invalid character '□' (U+2705)
```

Check Repayment Performance

Compare on-time payments between both groups.

```
df.groupby('risk_category')['On-time Payments (%)'].mean()
```

□ Big difference! Prime customers have a much higher on-time payment rate.

Compare Number of Loan Increases

More increases = more exposure to risk?

```
df.groupby('risk_category')['No. of Increases in 2023'].mean()
```

□ Very similar—both groups get almost the same number of limit increases.

Key Takeaways:

Near-Prime customers have much lower on-time payment rates (~87% vs. 97% for Prime).

1. This suggests higher risk of default in Near-Prime customers.
2. The loan amount & number of increases aren't the main problem—repayment behavior is!

Despite lower on-time payments, Near-Prime customers still receive similar loan increases.

1. This might indicate a flaw in the approval policy—it's not considering past repayment performance strongly enough.

Compare Default Risk & On-Time Payments using a Scatter Plot

```
import seaborn as sns
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(8,5))
sns.scatterplot(data=df, x='On-time Payments (%)', y='Estimated
Default Risk (%)', hue='risk_category')
plt.title("Relationship Between On-Time Payments & Default Risk")
plt.show()
```

Check Correlation Between Default Risk & Other Features

```
df.corr(numeric_only=True)['Estimated Default Risk
(%)'].sort_values(ascending=False)
```

Analysis of Correlation Results

Key Observations:

On-time Payments (%) has a perfect negative correlation (-1.000000) with Estimated Default Risk (%)

1. This means higher on-time payments → lower default risk (expected).
2. All other variables have weak or near-zero correlation with default risk.
3. Total Profit Contribution (\$) (0.0016) and No. of Increases in 2023 (0.0001) are almost uncorrelated with default risk.
4. Initial Loan (\$) (-0.0034) and Days Since Last Loan (-0.0019) have near-zero impact.

What This Means for Loan Limit Optimization:

5. On-time payments (%) should be the primary factor in determining loan increases.
6. Since "No. of Increases in 2023" and "Total Profit Contribution (\$)" do not correlate with default risk, we may need advanced models (Markov Chains, Simulation) instead of just correlation-based approaches.
7. External economic factors (inflation, employment rates) might be needed to refine risk assessment.

□ Next Step: Start Markov Chain Modeling

Step 1: Understanding the Markov Chain Modeling:

A Markov Chain is a stochastic model that predicts future states based on the current state. In our case, a borrower's risk category (Prime, Near-Prime, etc.) changes over time based on repayment behavior.

We'll define:

1. States → Risk categories: Prime, Near-Prime, Subprime (if applicable).
2. Transitions → Probability of moving from one state to another.
3. Transition Matrix → A matrix showing probabilities of movement between risk states.

Step 2: Building a Transition Matrix

```
df['Customer ID'].value_counts().head(10)
```

This means each Customer ID appears only once in the dataset, so we don't have historical transitions for each borrower.

Step 3 : Check Risk Category Distribution

```
df['risk_category'].value_counts(normalize=True)
```

step 4: Manually Define a Transition Matrix

```
import numpy as np
import pandas as pd

# Define states
states = ['Prime', 'Near-Prime']

# Assume some transition probabilities (we will refine this later)
transition_matrix = pd.DataFrame(
    np.array([[0.85, 0.15], # Prime → (Prime, Near-Prime)
              [0.20, 0.80]]), # Near-Prime → (Prime, Near-Prime)
    index=states,
    columns=states
)

transition_matrix
```

Step 5: Applying the Markov Model

Define the Initial Distribution

```
initial_distribution =
df['risk_category'].value_counts(normalize=True).sort_index()
initial_distribution
```

Predict Future Risk Category Distribution

```
# Convert to numpy array
current_state = initial_distribution.values

# Predict risk distribution after N steps (e.g., 6 months)
n_steps = 6
for _ in range(n_steps):
    current_state = np.dot(current_state, transition_matrix.values)

predicted_distribution = pd.Series(current_state,
index=transition_matrix.columns)
predicted_distribution
```

1. **1** Good news: More customers are becoming Prime → we can safely offer them better loan limits.
2. **2** Caution: 41% still in Near-Prime → we need risk controls before increasing their limits.

step 6: Simulating Loan Limits for Different Risk Groups

Define Loan Limit Increase Probabilities

```
loan_acceptance_prob = {
    "Prime": 0.75, # 75% of Prime borrowers accept increases
    "Near-Prime": 0.50 # 50% of Near-Prime borrowers accept increases
}

repayment_prob = {
    "Prime": 0.90, # 90% of Prime borrowers repay on time
    "Near-Prime": 0.70 # 70% of Near-Prime borrowers repay on time
}
```

Simulate Loan Limit Increases

```
df['loan_increase_acceptance'] =
df['risk_category'].map(loan_acceptance_prob)
df['successful_repayment'] = df['risk_category'].map(repayment_prob)
df[['risk_category', 'loan_increase_acceptance',
'successful_repayment']].head()
```

Results confirm our assumptions:

1. Prime borrowers (0.75 acceptance, 0.9 repayment) → More reliable, accept loan increases more often, and repay successfully.

2. Near-Prime borrowers (0.50 acceptance, 0.7 repayment) → Riskier, accept fewer increases, and have a lower repayment rate.

Step 7: Expected Profit Calculation

Define Expected Profit Formula

Expected Profit=(Acceptance Rate)×(Repayment Probability)×40

```
df['expected_profit'] = df['loan_increase_acceptance'] *  
df['successful_repayment'] * 40  
df[['risk_category', 'expected_profit']].head()
```

step 8: Annualized Profit Estimation

Since customers can get up to 6 increases per year, let's estimate their annual expected profit:

Formula for Annual Profit

Annual Profit=Expected Profit×min(6,No. of Increases in 2023)

```
df['annual_expected_profit'] = df['expected_profit'] * df['No. of  
Increases in 2023'].clip(upper=6)  
df[['risk_category',  
'annual_expected_profit']].groupby('risk_category').mean()
```

Prime borrowers generate nearly 2x the profit compared to Near-Prime!

Step 9: Risk-Adjusted Profit Calculation

Formula for Risk-Adjusted Profit

$\text{Risk-Adjusted Profit} = \text{Annual Profit} \times (1 - \text{Estimated Default Risk (\%)})$

```
df['risk_adjusted_profit'] = df['annual_expected_profit'] * (1 -  
df['Estimated Default Risk (%)'] / 100)  
df[['risk_category',  
'risk_adjusted_profit']].groupby('risk_category').mean()
```

Near-Prime borrowers lose \$4 due to higher defaults, while Prime borrowers only lose \$1.

Step 10: Optimizing Loan Limit Increase Strategy

Define Decision Rule

We need to approve increases only if:

Risk-Adjusted Profit>Minimum Required Profit Threshold

1. Should we set the threshold at \$30 to ensure profitability?
2. Or should we allow some Near-Prime borrowers if they improve over time?

```
profit_threshold = 30
df['approved_for_increase'] = df['risk_adjusted_profit'] >
profit_threshold
df['approved_for_increase'].value_counts()
```

Insights from Approval Results

1. 16,793 customers approved for a loan limit increase.
2. 13,207 customers not approved based on our risk-adjusted profit threshold.

Step 11: Evaluating the Policy

1. Approval Rate: What % of total customers got approved?

```
approval_rate = df['approved_for_increase'].mean() * 100
print(f"Approval Rate: {approval_rate:.2f}%")
```

2. Approval by Risk Category

```
print(df.groupby('risk_category')
      ['approved_for_increase'].value_counts(normalize=True) * 100)
```

Step: 12 Risk vs. Profit Tradeoff

1. Average profit contribution of approved vs. rejected customers
2. Default risk distribution in approved cases

```
df.groupby('approved_for_increase')[['Total Profit Contribution ($)',
                                     'Estimated Default Risk (%)']].mean()
```

Key Takeaways

Profitability is only generated from approved customers (expected).

Default risk is nearly the same (9.99% vs. 9.97%)

1. This suggests rejections aren't significantly reducing default risk.
2. We may be rejecting customers unnecessarily, leading to lost revenue.

step 13: profitability by risk category

```
df[df['approved_for_increase'] == True].groupby('risk_category')
['Total Profit Contribution ($)'].mean()
```

Key Takeaways

Both Prime & Near-Prime customers contribute nearly the same profit!

This suggests Near-Prime approvals are just as profitable as Prime.

1. We might be over-restricting Near-Prime customers without much financial gain.
2. Since risk was nearly the same (9.97% vs. 9.99%), we should explore if we can approve more Near-Prime customers without significantly increasing default risk.

Step 14: Approval rates for different risk categories:

```
df.groupby('risk_category')['approved_for_increase'].mean() * 100
```

Insights

Approval rates are nearly identical for both risk categories.

1. This suggests the model is not overly favoring Prime customers.
2. However, a slightly higher approval rate for Prime aligns with traditional lending strategies.

Near-Prime is being approved at a competitive rate (55.85%), but we should check if their default rate justifies this.

3. If Near-Prime defaults significantly more, the approval rate might still be too high.
4. If default rates are similar, we could consider raising Near-Prime approvals to capture more profit.

step 15: Default rates by risk category

```
df.groupby('risk_category')['Estimated Default Risk (%)'].mean()
```

Key Observations

Near-Prime defaults nearly 5x more than Prime customers.

1. This is a significant risk factor.

1. Despite a 55.85% approval rate, Near-Prime customers default much more, potentially hurting profitability.

Prime customers are significantly safer.

2. Their 2.55% default rate makes them much more reliable.
3. This justifies why they should receive more approvals or higher loan limits.

Step 15: Approved Near-Prime customers have a lower default rate

```
df.groupby(['approved_for_increase', 'risk_category'])['Estimated Default Risk (%)'].mean()
```

Step 16: Improve approval decisions,

Analyze more factors influencing Near-Prime defaults (e.g., payment history, loan size). Refine approval thresholds (e.g., only approve Near-Prime customers with 80%+ on-time payments).

Next Steps: Optimizing Loan Approval Decision

Step 1: Define an Approval Model

We can use one of these approaches: 1 Rule-based model (e.g., approve if on-time payment > X)
2 Logistic regression (predict approval likelihood) 3 Machine learning models (Random Forest, XGBoost, etc.) 4 Reinforcement learning (dynamic decision-making)

□ Step 1: Train Logistic Regression Model for Loan Approval

We'll predict approved_for_increase based on:

1. On-time Payments (%)
2. Estimated Default Risk (%)
3. No. of Increases in 2023
4. Risk Category (encoded as 0 = Near-Prime, 1 = Prime)

Step 2: Implement Logistic Regression

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report

# Selecting features and target variable
features = ['On-time Payments (%)', 'Estimated Default Risk (%)', 'No. of Increases in 2023']
X = df[features]
X['Risk Category'] = df['risk_category'].map({'Near-Prime': 0, 'Prime': 1}) # Encoding
y = df['approved_for_increase']
```

```

# Splitting the data
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Scaling features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Train Logistic Regression Model
model = LogisticRegression()
model.fit(X_train_scaled, y_train)

# Predictions
y_pred = model.predict(X_test_scaled)

# Evaluate Model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
print(classification_report(y_test, y_pred))

df_encoded = df.copy()

# Convert 'risk_category' to numeric values
df_encoded['risk_category'] =
df_encoded['risk_category'].map({'Prime': 1, 'Near-Prime': 0})

# Now plot the heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(df_encoded.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Feature Correlation Heatmap")
plt.show()

```

Perfect Accuracy (1.0)? That suggests:

1. Data Leakage – The model might be learning something too obvious from the dataset.
2. Overfitting – The test set might not be challenging enough.

Next steps Saving & Retaining the Model:

1. Save the Model:

```

import joblib

# Save the trained model

```

```
joblib.dump(model, "loan_approval_model.pkl")  
print("Model saved successfully!")
```

2. Load the Model (When Needed):

```
# Load the saved model  
model_loaded = joblib.load("loan_approval_model.pkl")  
  
# Now you can use model_loaded for predictions
```

3. Test the Loaded Model

```
sample_input = X_test.iloc[0:1] # Select one test sample  
prediction = model_loaded.predict(sample_input)  
print("Prediction:", prediction)
```

Here's what we covered:

1. ☐ Data Preprocessing & Analysis (handling missing values, feature engineering)
2. ☐ Risk Categorization & Customer Behavior Analysis
3. ☐ Approval Decision Modeling (classification model for approvals)
4. ☐ Evaluation of Approval Impact (profit, risk, approval rates)
5. ☐ Model Training & Performance Evaluation (precision, recall, F1-score)
6. ☐ Saving & Retaining the Model (for future use)

Everything is now ready for you.